

Chapter 7

Evaluating the Instruction Set
Architecture of H1: Part 1

We will study the assembly code generated by a “dumb” C++ compiler

- Will provide a sense of what constitutes a good architecture.
- Will provide a better understanding of high-level languages, such as C++ and Java.

Dumb compiler

- Translates each statement in isolation of the statements that precede and follow it.
- Dumb compiler lacks all non-essential capabilities, except for a few that would *minimally* impact its complexity.

Dump compiler translates

`x = 2;`

`y = x; to`

`ldc 2`

`st x`

`ld x ; this ld is unnecessary`

`st y`

Compiler does not remember what it did for the first instruction when it translates the second.

A smart compiler would translate

`x = 2;`

`y = x;` to

`ldc 2`

`st x`

`st y`

Smart compiler is able to generate more efficient code for the second instruction by remembering what it did for the first instruction.

Dumb compiler generates code left to right

```
d = a + b + c;
```

where **a**, **b**, **c**, and **d** are global variables, is translated to

```
ld    a           ; access a first
add   b           ; access b next
add   c           ; access c last
st    d
```

Dumb compiler follows order required by operator precedence and associativity.

* has higher precedence than +
= is right associative.

```
d = a + b*c;    // code for b* c first
```

```
a = b = c = 5; // code for c = 5 first
```

When dumb compiler generates code for a binary operator, it accesses operands in the order that yields better code.

```
a = b + 1;
```

our compiler can generate an `ldc` instruction to get the constant 1, but only if it generates this code before it generates code to access `b`:

```
ldc  1
add  b
st   a
```

If `b` were accessed first, then the `ldc` instruction could not be used to load 1 (it would destroy the value of `b` in the `ac` register). Thus, the compiler would have to generate

```
ld   b
add  @1
st   a
```

where `@1` is a compiler-generated label defined as

```
@1:  dw    1
```


Dumb compiler always follows
order imposed by parentheses

$v = (w + x) + (y + z);$

1st

3rd

2nd

Other orders might yield more
efficient code.

Dumb compiler performs constant folding (compile-time evaluation of constant expressions)

`x = 2 + 3;`

a compiler can generate

```
ldc  2
add  @3
st   x
```

Run-time evaluation

where `@3` is defined as

```
@3:  dw    3
```

Thus, at run time the `add` instruction computes the value of `2 + 3`. Alternatively, a compiler can add 2 and 3 at compile time to get 5, and then generate code to load and store 5:

```
ldc  5
st   x
```

Compile-time evaluation

Constant folding only on
constants at the *beginning* of an
expression.

$y = 1 + 2 + x;$ // constant folding

$y = 1 + x + 2;$ // no constant folding

$y = x + 1 + 2;$ // no constant folding

Global variables in C++

- Declared outside of a function definition.
- Scope is from point of declaration to end of file (excluding functions with identically named local variable).
- Default to an initial value of 0 if an initial value is not explicitly specified.
- Are translated to a `dw` statement in assembler code.

FIGURE 7.1

```
1 #include <iostream>
2 using namespace std;
3
4 int x;           // global variable with no initial value specified
5 void fa()
6 {               // cannot reference y from within fa
7     x = x + x + 2;
8 }
9 int y = 3;      // global variable with initial value specified
10 void fb()
11 {
12     x = 5;
13     y = -2;
14     fa();
15 }
16 void main()
17 {
18     fb();
19     cout << "x = " << x << endl;    // displays "x = 12"
20     cout << "y = " << y << endl;    // displays "y = -2"
21 }
```

x and y in preceding slide are
translated to

x: dw 0 ; 0 default value

y: dw 3

Direct instructions are used to
access global variables

`x = 5;`

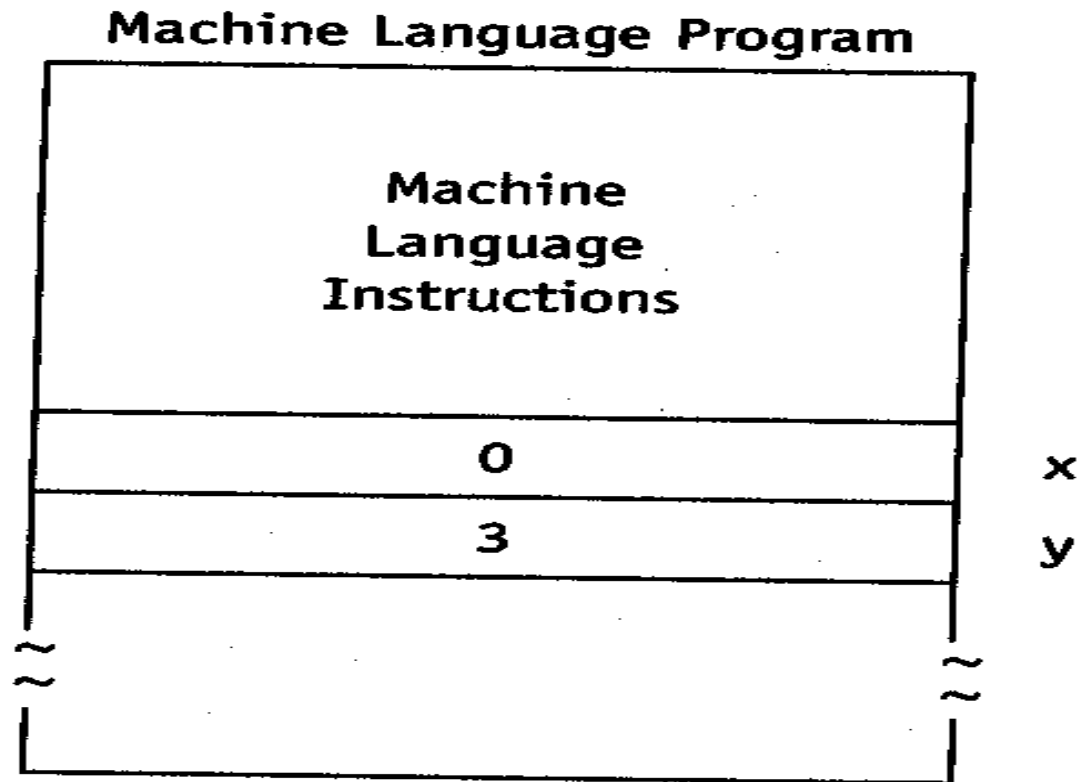
is translated to

`ldc 5`

`st x ; use direct instruction`

Global variables are part of the machine language program

FIGURE 7.2



Compiler generated names start with '@'

@2 is the label generated for the constant 2

@_2 is the label generated for the constant -2

@m0, @m1, ... , are for string constants

See @2, @_2, @m0, @m1 in the next slide.

FIGURE 7.3

```

1 fa:      ld      x      ; x = x + x + 2;
2          add     x
3          add     @2
4          st      x
5
6          ret
7 ;=====
8 fb:      ldc     5      ; x = 5;
9          st      x
10
11         ld      @_2     ; y = -2;
12         st      y
13
14         call    fa      ; fa()
15
16         ret
17 ;=====
18 main:    call    fb      ; fb()
19
20         ldc     @m0     ; cout << "x = " << x << endl;
21         sout
22         ld      x
23         dout
24         ldc     '\n'
25         aout
26
27         ldc     @m1     ; cout << "y = " << y << endl;
28         sout
29         ld      y
30         dout
31         ldc     '\n'
32         aout
33
34         halt
35 x:       dw      0      ; global variable
36 y:       dw      3      ; global variable
37 @2:      dw      2      ; @2 is a compiler-generated name
38 @_2:     dw      -2     ; @_2 is a compiler-generated name
39 @m0:     dw      "x = "  ; @m0 is a compiler-generated name
40 @m1:     dw      "y = "  ; @m1 is a compiler-generated name
41         end main

```

A local variable can be referenced by name only in the function or sub-block within the function in which it is defined.

- Two types of local variables: *dynamic* and *static*.
- Static local variables are defined with a `dw` statement. Direct instructions are used to access them. Created at *assembly time*.
- Dynamic local variables are created on the stack at *run time*. Relative instructions are used to access them.

FIGURE 7.4

```
1 #include <iostream>
2 using namespace std;
3
4 void fc()
5 {
6     int dla ;           // dynamic local variable
7     int dlb = 7;        // dynamic local variable
8
9     static int sla;      // static local variable
10    static int slb = 5;   // static local variable
11
12    cout << sla << endl;
13    cout << slb << endl;
14    dla = 25;
15    dlb = 26;
16    sla = 27;
17    slb = 28;
18
19 }
20 void main()
21 {
22     fc();
23     // sla and slb retain values 27 and 28 between calls
24     fc();
25 }
```

Output from preceding program

```
0      (initial value of sla on first call of fc)
5      (initial value of slb on first call of fc)
27     (initial value of sla on second call of fc)
28     (initial value of slb on second call of fc)
```

sla, slb translated to

```
@s0_sla:      dw      0
@s1_slb:      dw      5
```

Thus, the C++ statements

```
sla = 27;
slb = 28;
```

are translated to

```
ldc    27
st      @s0_sla
ldc    28
st      @s1_slb
```

Why are the names of static local variables not carried over *as is* (as with global variables) to assembler code?

See the next slide.

FIGURE 7.5

	C++ Program	Assembler Program
1		
2		
3	int x;	
4	void f1()	f1:
5	{	.
6	static int x;	.
7	.	.
8	.	.
9	.	ret.
10	}	
11	void f2()	f2:
12	{	.
13	static int x;	.
14	.	.
15	.	.
16	.	ret
17	}	
18	void main()	main:
19	{	.
20	.	.
21	.	.
22	.	halt
23	}	.
24		.
25		x: dw 0 ; global
26		x: dw 0 ; static local in f1
27		x: dw 0 ; static local in f2

↑
Error: Duplicate labels


A dynamic local variable is allocated on the stack with a push or aloc instruction.

push is used if the variable has to be initialized. aloc is used otherwise.

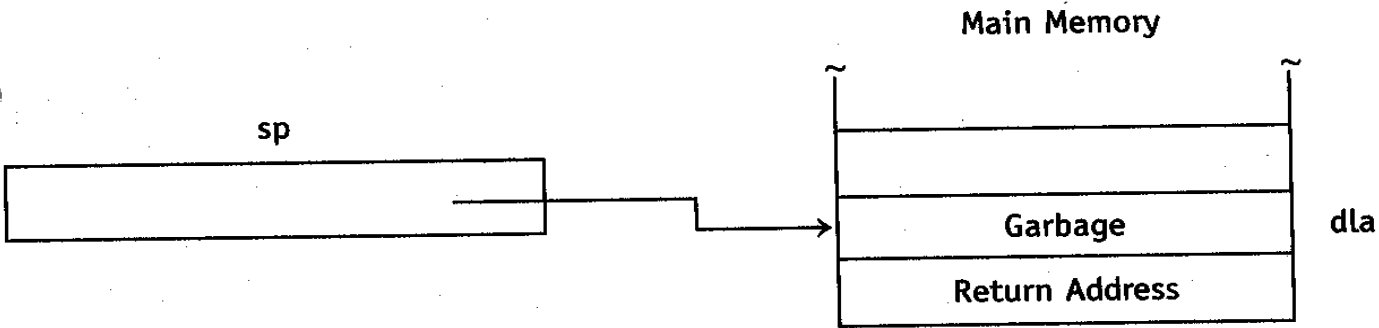
FIGURE 7.7

```
1      fc:          aloc 1          ; int dla;
2
3          ldc 7          ; int dlb = 7;
4      push
5          .
6          .
7          .
8          dloc 2          ; deallocate dla and dlb
9      ret
10     ;=====
11
12     main:        call fc          ; fc();
13
14          call fc          ; fc();
15
16          halt
17     @s0_sla:     dw 0          ; static local variable sla
18     @s1_slb:     dw 5          ; static local variable slb
19     end main
```


FIGURE 7.6 a) On entry into the function ϵc



b) After the creation of $d1a$



c) After the creation of $d1b$



The relative addresses of dla and
dlb are 1 and 0, respectively.

See the preceding slide.

FIGURE 7.8

```

1 fc:      aloc 1          ; int dla;
2
3          ldc 7           ; int dlb = 7;
4          push
5
6          ld @s0_sla      ; cout << sla << endl;
7          dout
8          ldc '\n'
9          aout
10
11         ld @s1_slb      ; cout << slb << endl;
12         dout
13         ldc '\n'
14         aout
15
16         ldc 25          ; dla = 25;
17         str 1
18
19         ldc 26          ; dlb = 26;
20         str 0
21
22         ldc 27          ; sla = 27;
23         st @s0_sla
24
25         ldc 28          ; sla = 28;
26         st @s1_slb
27
28         dloc 2          ; remove dla and dlb
29         ret
30 ;=====
31 main:    call fc        ; fc();
32
33         call fc        ; fc();
34
35         halt
36 @s0_sla:  dw 0           ; static local variable sla
37 @s1_slb:  dw 5           ; static local variable slb
38         end main

```

The default initialization of global and static local variables (to 0) does not require additional time or space. But the default initialization of dynamic local variables would require extra time and space (a `ld` or `ldc` followed by a push versus a single `aloc`).

x and y initially 0; z initially has garbage

FIGURE 7.9

```
1 int x;                // global, initial value = 0
2 void fd()
3 {
4     static int y;      // static local, initial value = 0
5     int z;             // dynamic local, initial value undefined
6     .
7     .
8     .
9 }
10 void main()
11 {
12     fd();
13 }
```

Default initializations are often not used, in which case the extra time and space required for the initialization of dynamic local variables would be wasted.

FIGURE 7.10

```
void fe() {  
    int z;  
    .  
    .  
    z = x + y;  
    .  
    .  
}
```


Relative addresses can change during the execution of a function

- Changing relative addresses makes compiling a more difficult task—the compiler must keep track of the correct relative address.
- Changing relative addresses is an frequent source of bugs in hand-written assembly code for the H1.

Relative address of x changes

FIGURE 7.11

C++	Assembly Language
1 void legal()	legal:
2 {	
3 int x;	aloc 1 ; allocate x
4	
5 x = 3;	ldc 3
6	str 0 ; relative address of x is 0
7	
8 int y;	aloc 1 ; allocate y--changes rel add of x
9	
10 x = 5;	ldc 5
11	str 1 ; relative address of x is 1
12	
13	dloc 2
14 }	ret

Call by value

- The value (as opposed to the address) of the argument is passed to a function.
- An argument and its corresponding parameter are distinct variables, occupying different memory locations.
- Parameters are created by the *calling function* on function call; destroyed by the *calling function* on function return.

m is the argument; x is its parameter; y is a local variable

FIGURE 7.12

```
1 int m = 5;           // m is a global variable
2 void fg(int x)       // x is the parameter
3 {
4     int y;           // y is a dynamic local variable
5     y = 7;
6     x = y + 2;
7 }
8 void main()
9 {
10     fg(m);          // m is the argument
11 }
```

Creating parameters

The parameter x comes into existence during the call of fg . x is created on the stack by pushing the value of its corresponding argument (m).



Start of function call f(m)

Push value or argument, thereby creating the parameter.

ld m

push ; creates x

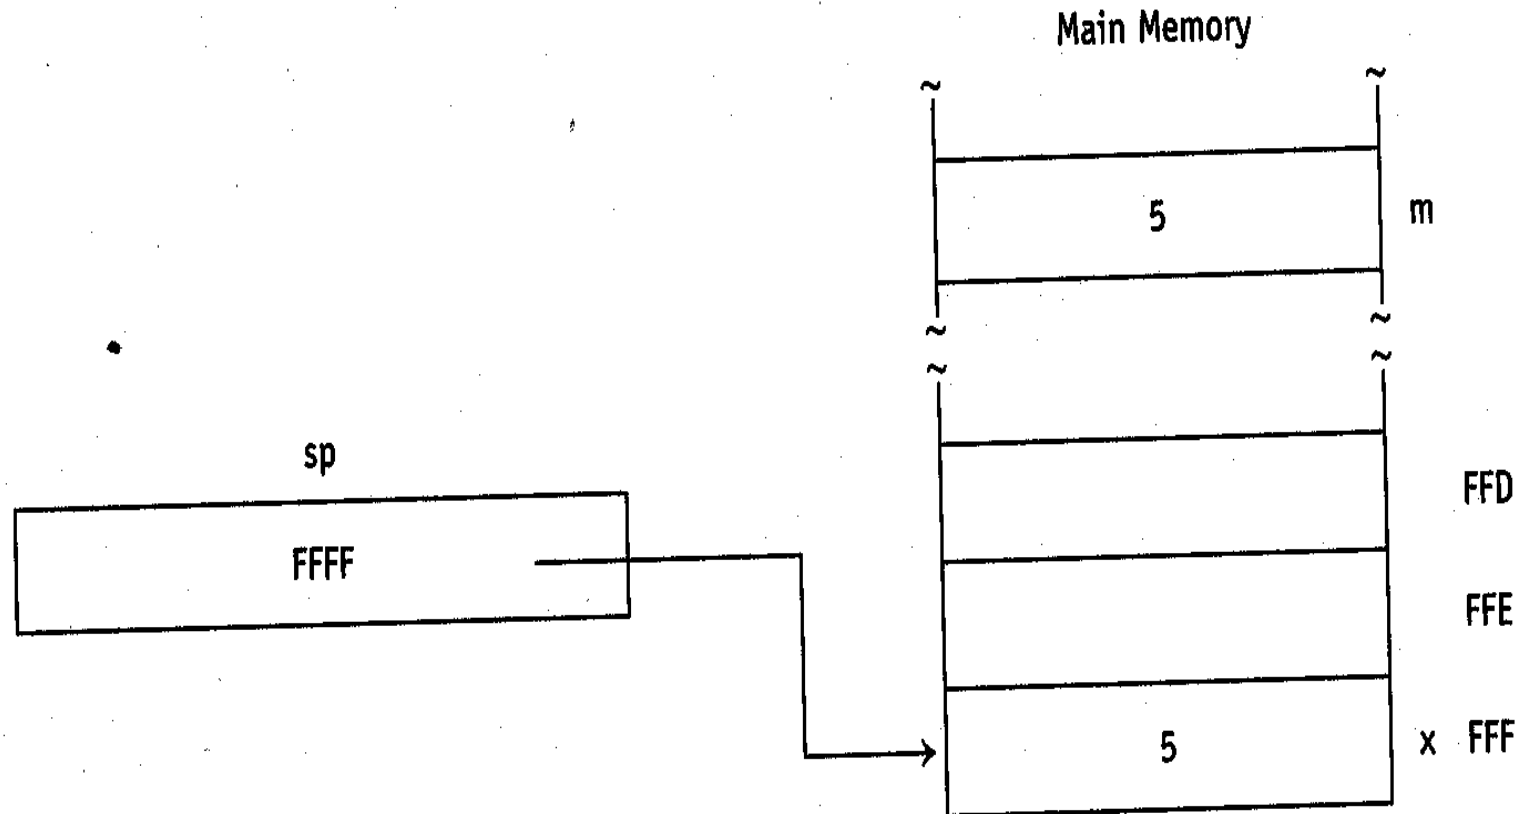
FIGURE 7.12A



Call the function (which pushes return address on the stack)

call fg


FIGURE 7.12B



Allocate uninitialized local variable with alloc

alloc 1

FIGURE 7.12C



Execute function body

ldc 7 ; y = 7;


str 0

ldc 2 ; x = y + 2;

addr 0

str 2

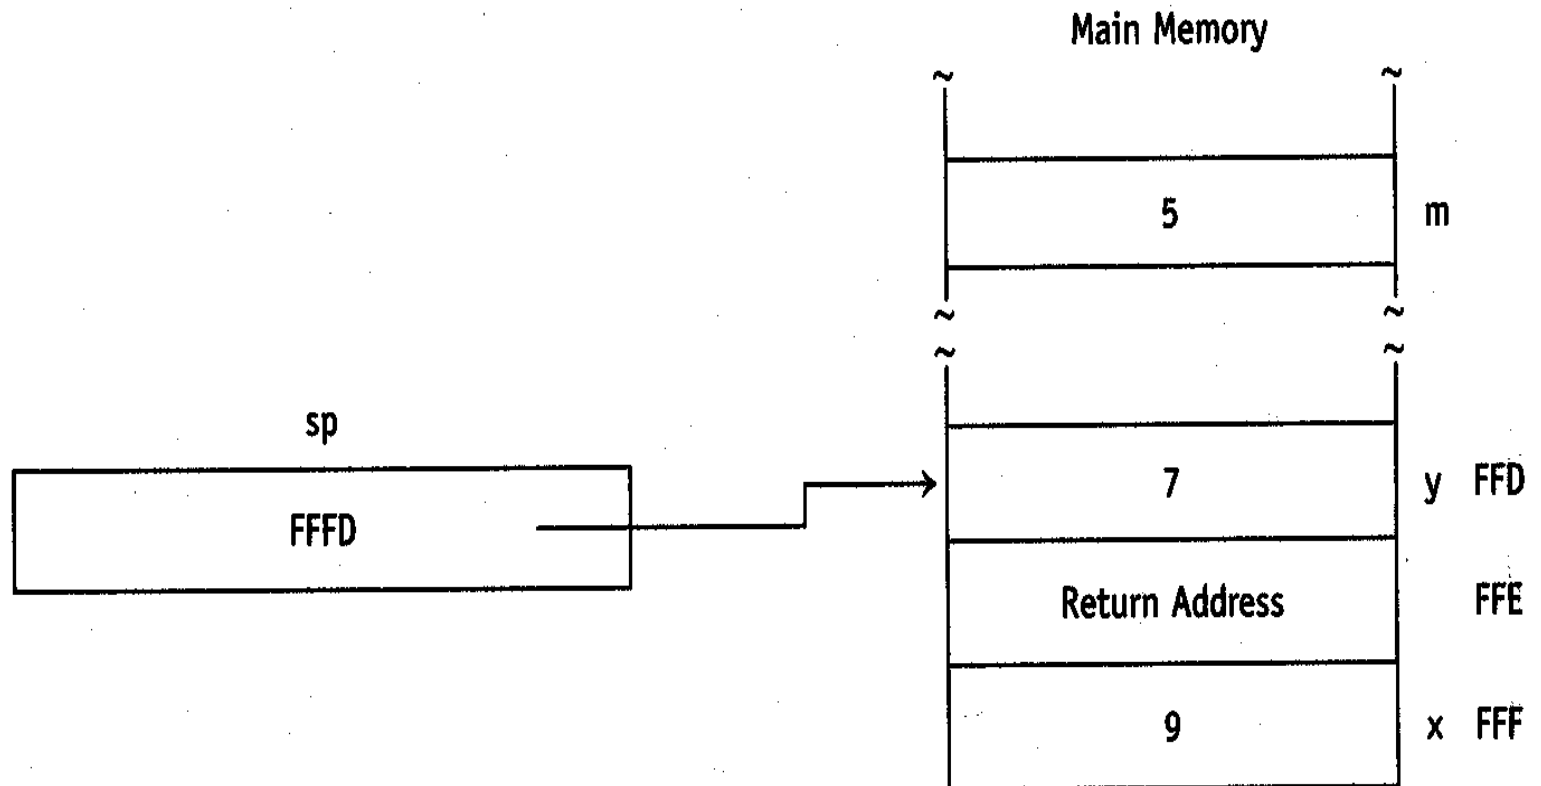
FIGURE 7.12D



Deallocate dynamic local variable

dloc 1

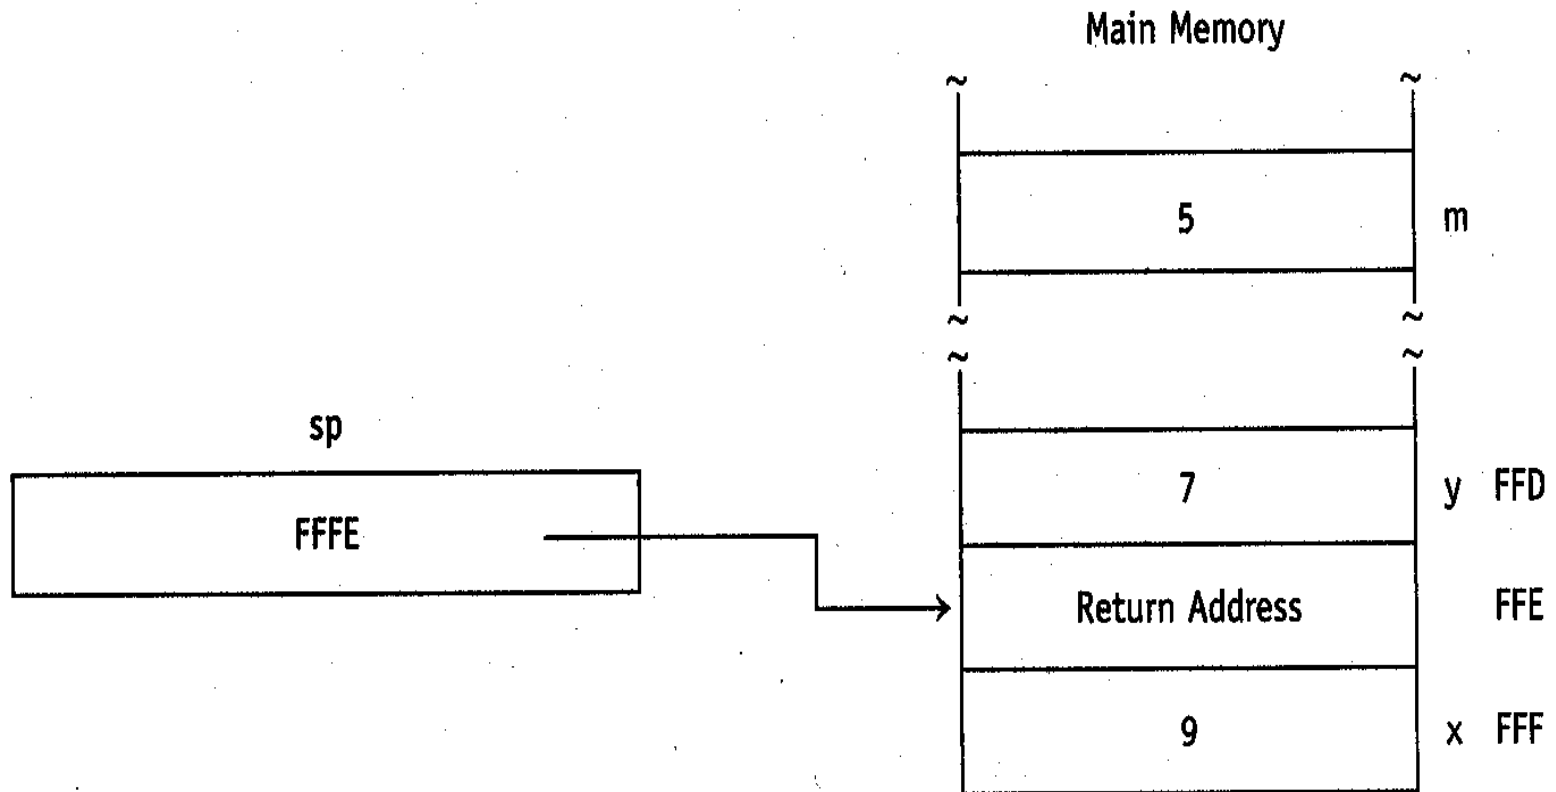
FIGURE 7.12E



Return to calling function (which pops the return address)

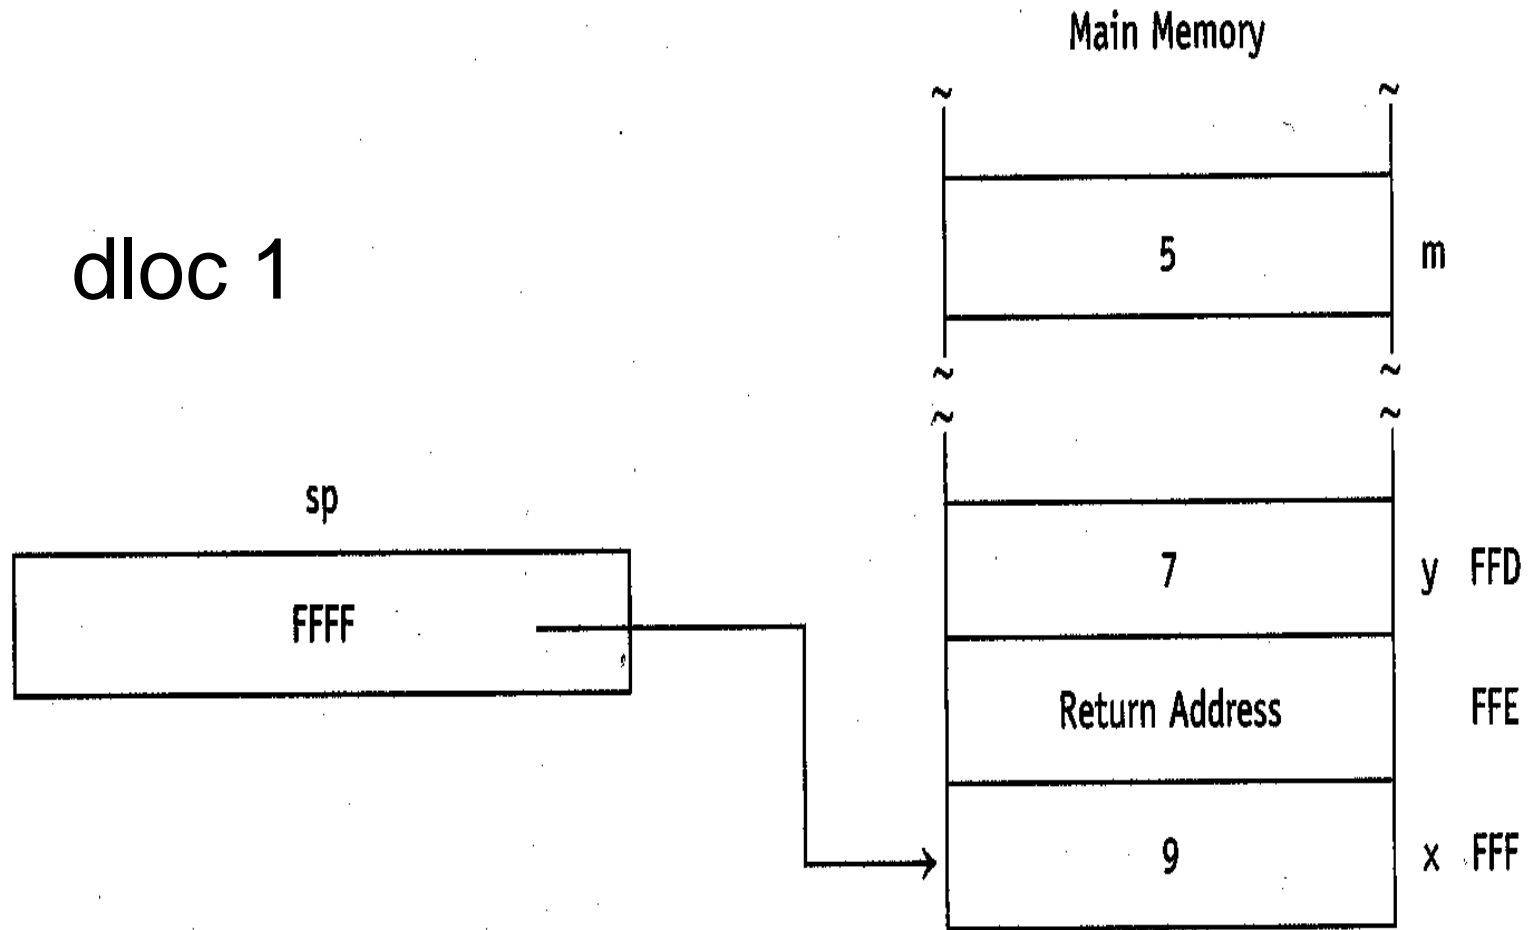
ret

FIGURE 7.12F



Deallocate parameter

FIGURE 7.12G



Back to initial configuration

FIGURE 7.12H





FIGURE 7.13

```
1 fg:          aloc 1                ; int y;
2
3             ldc 7                  ; y = 7;
4             str 0
5
6             ldc 2                  ; x = y + 2;
7             addr 0
8             str 2
9
10            dloc 1                ; deallocate y
11            ret
12 ;=====
13 main:        ld m                  ; fg(m);
14             push
15             call fg
16             dloc 1
17
18             halt
19 m:           dw 5
26            end main
```

Parameter creation/destruction

The *calling function* creates and destroys parameters.

```
ld      m  
push    ; create parameter x  
call    fg  
dloc    1 ; destroy parameter x
```

Dynamic local variable creation/destruction

The *called function* creates and destroys dynamic local variables.

alloc 1 ; create y.

▪

▪

dloc 1 ; destroy y

Push arguments right to left

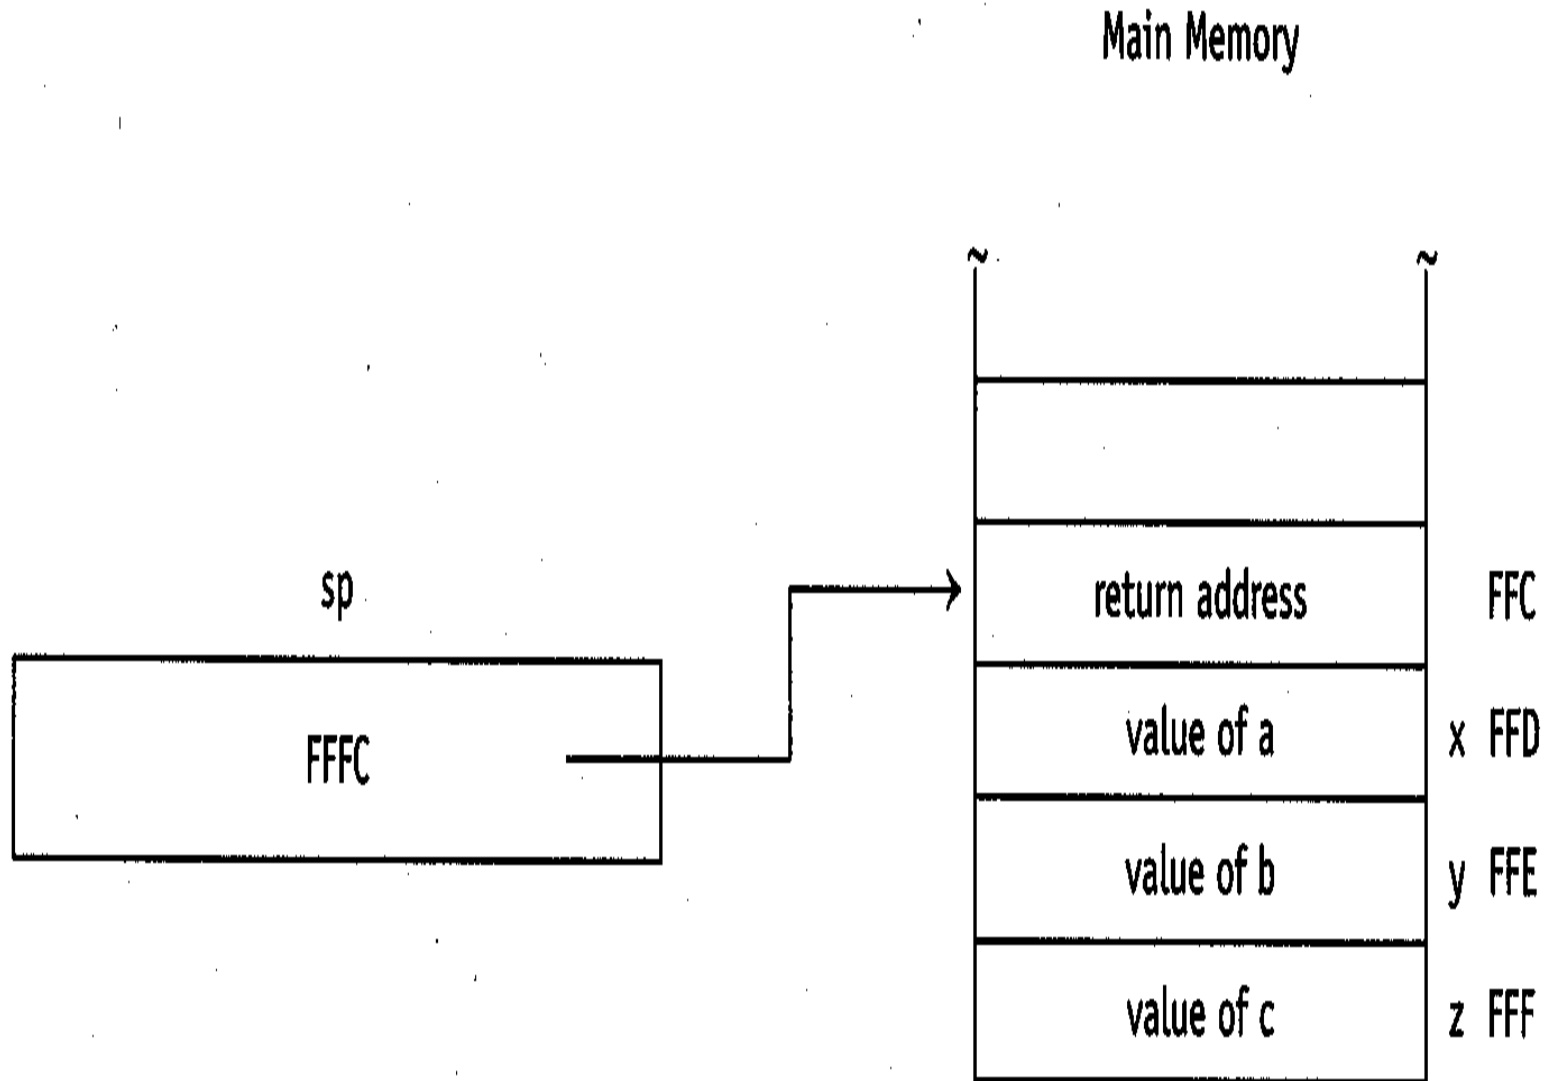
```
void fh(int x, int y, int z)
{
    ...
}
```

If **a**, **b**, and **c** are global variables, then the call,
`fh(a, b, c);`
is translated to

```
ld    c
push           ; create z
ld    b
push           ; create y
ld    a
push           ; create x
call  fh
dloc 3         ; remove all three parameters
```

On entry into fh
x, y, and z have relative addresses 1, 2, and 3

FIGURE 7.14



The C++ return statement returns values via the ac register.

FIGURE 7.15

```
1 #include <iostream>
2 using namespace std;
3
4 int x, y;
5 int add_one(int z)
6 {
7     return z + 1;
8 }
9 void main()
10 {
11     x = 1;
12     y = add_one(x);    // value returned is assigned to y
13     cout << "y = " << y << endl;
14 }
```

Using the value returned in the ac

```
y = add_one(x);
```

is

```
ld    x        ; get value of x
push           ; create parameter on the stack
call add_one
dloc 1          ; remove parameter
st     y        ; use value returned in ac
```

FIGURE 7.16

```
1 add_one:    ldc  1          ;return z + 1;
2            addr 1
3            ret
4 ;=====
5 main:       ldc  1          ; x = 1;
9            st   x
10
11           ld   x          ; y = add_one(x);
12           push
13           call add_one
14           dloc 1
15           st   y
16
17           ldc  @m0        ; cout << "y = " << y << endl;
18           sout
19           ld   y
20           dout
21           ldc  '\n'
22           aout
23
24           halt
25 x:         dw   0
26 y:         dw   0
27 @m0:       dw   "y = "
28           end   main
```

Is it possible to replace relative instructions with direct instructions?

Answer: sometimes

It *is* possible to replace relative instructions accessing x, y, z with direct instructions in this program.


FIGURE 7.17

```
1 void sum(int x, int y)
2 {
3     int z;
4
5     z = x + y;
6 }
7 void main()
8 {
9     sum(5, 10);
10 }
```


Can use either an absolute address or its corresponding relative address

	relative address	absolute address
x	2	FFE
y	3	FFF
z	0	FFC

FIGURE 7.18



Stack frame: the collection of parameters, return address and local variables allocated on the stack during a function call and execution.

A stack frame is sometimes called an *activation record*.

FIGURE 7.19

```
1 sum:      aloc 1
2
3          ldr  2      ; replace with ld  FFEh
4          addr 3      ; replace with add FFFh
5          str  0      ; replace with st  FFCh
6
7          dloc 1
8          ret
9  ; =====
10 main:    ldc  10      ; sum(5,10);
11
12          push
13
14          ldc  5
15
16          push
17
18          call sum
19
20          dloc 2
21
22          halt
23
24          end  main
```

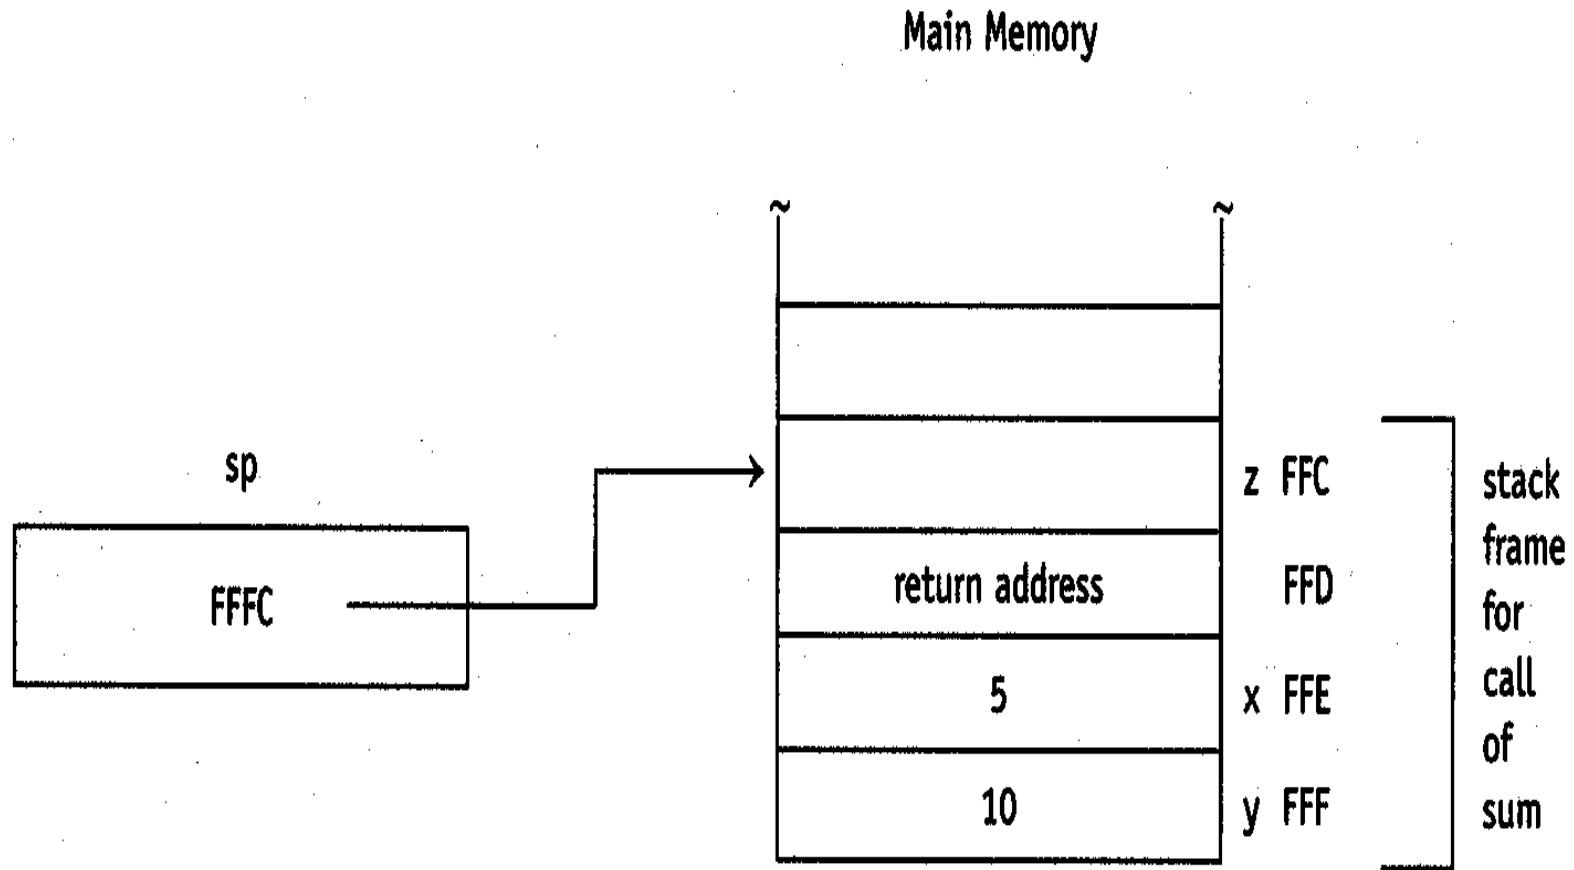
But for the following program, the *absolute addresses* of x, y, and z are different during the two calls of **sum**. But the *relative addresses* are the same.

FIGURE 7.20

```
1 void sum(int x, int y)
2 {
3     int z;
4
5     z = x + y;
6 }
7 void fk()
8 {
9     sum(5, 10);
10 }
11 void main()
12 {
13     sum(5, 10);
14     fk();
15 }
```

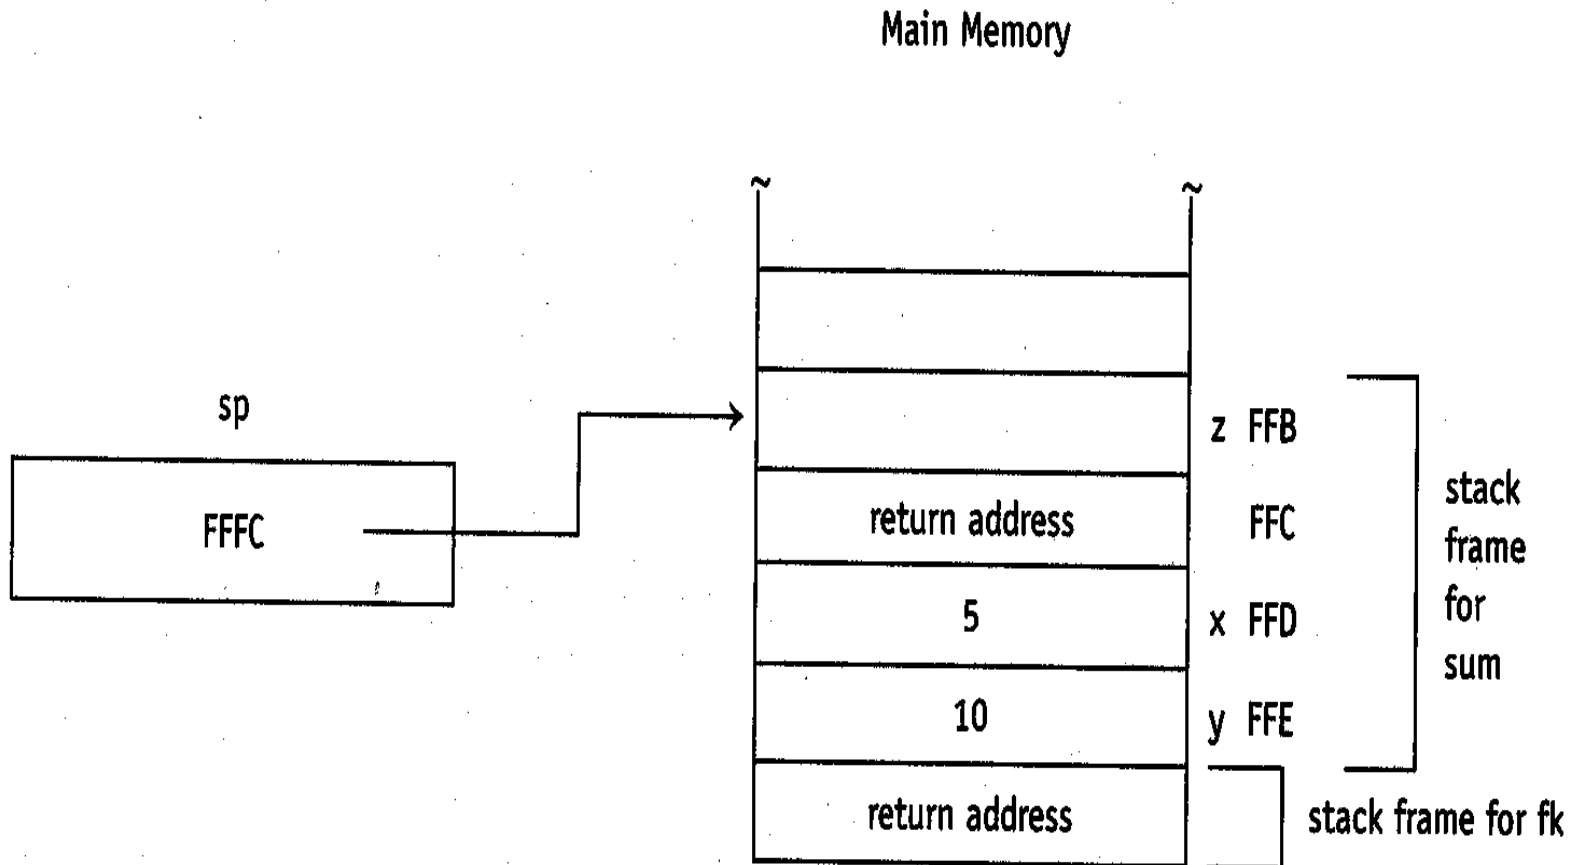
Stack frame for 1st call of sum

FIGURE 7.18



Stack frame for 2nd call of sum

FIGURE 7.21



Another problem with H1: determining the absolute address of items on the stack.

Determining the addresses of globals and static locals is easy, but difficult for dynamic locals.

Easy

Suppose `x` and `y` are global variables. Because `x` and `y` are defined with `dw` directives, the C++ statement

```
y = &x;    // "&" means "address of"
```

is translated to the assembly code

```
ldc  x      ; load address of x  
st   y      ; store in y
```

Also easy

time. Now suppose **x** and **y** are static local variables. Then

y = &x;

would be translated to

```
ldc  @s0_x
```

```
st   @s1_y
```


Relative address of x below is 1. What is its absolute address? It varies! Thus, &x must be computed at *run time* by converting its relative address to its corresponding absolute address.

FIGURE 7.22


```
1 void fm()
2 {
3     int x;
4     int *p;
5
6     p = &x;           // assigning address of local variable
7 }
8 void fn()
9 {
10    fm();
11 }
12 void main()
13 {
14    fm();
15    fn();
16 }
```

Absolute address of x different for 1st and 2nd calls of fm but relative address is the same (1).

FIGURE 7.23 a)



b)



To convert relative address 1 to absolute address:

```
    swap                ; corrupts sp
    st @spsave
    swap                ; restores sp
    ldc 1
    add @spsave
```

where

```
@spsave:    dw    0
```

FIGURE 7.24

```
1 fm:      alloc 1      ; int x;
2
3          alloc 1      ; int *p;
4
5          ; p = &x;
6          swap          ; get sp
7          st    @spsave ; save it
8          swap          ; restore sp
9          ldc   1       ; get relative address of x (1)
10         add   @spsave  ; convert to an absolute address
11         str   0        ; store absolute address in p
12
13         dloc  2        ; deallocate x and p
14         ret
15 ;      =====
16 fn:     call fm        ; fm();
17         ret
18 ;      =====
19 main:   call fm        ; fm();
20         call fn        ; fn();
21         halt
22 @spsave: dw    0
23         end   main
```

It is dangerous to corrupt the sp register on most computers, even for a short period of time.

An *interrupt mechanism* uses the stack. Interrupts can occur at any time. If the sp is in a corrupted state when an interrupt occurs, a system failure will result.

H1 does *not* have an interrupt mechanism.

Dereferencing pointers

Code generated depends on the type of the variable pointed to.

That is why you have to specify the type of the pointed-to item when declaring a pointer.


```
int    *p;    // p points to one-word number
long   *q;    // q points to two-word number
```

```
*p = 0;
```

zeros one word, the `int` item pointed to by `p`. But

```
*q = 0;
```

zeros two words, the `long` item pointed to by `q`. For the former statement the compiler generates

```
ldc    0
push
ld      p      ; get address in p
sti
```

But for the latter statement the compiler has to generate a completely different sequence:

```
ldc    0      ; prepare for sti
push
ld      q      ; get address in q
sti      ; put zero in first word
ldc    0
push      ; prepare for sti
ldc    1      ; get 1
add     q      ; get (address in q) + 1
sti      ; put zero in second word
```