

Chapter 16

Java Virtual Machine

To compile a java program in Simple.java, enter

```
javac Simple.java
```

javac outputs Simple.class, a file that contains *bytecode* (machine language for the *Java Virtual Machine* (JVM)).

To run Simple.class, enter

```
java Simple
```

java (the Java interpreter) makes
your computer act like the JVM
so it can execute bytecode.

Why is Java slow?

- Interpretation of bytecode can involve a lot of overhead.
- JVM dynamically links classes.
- JVM performs checks during loading, linking, and executing bytecode.

Why is Java good for the Web?

- Bytecode is space efficient.
- Bytecode is portable to any system with a java interpreter.
- Java applets are safe to run.

Four parts of the JVM

- Execution engine (contains pc register)
- Method area (contains information on each class: bytecode, static variables, information needed for verification and linking).
- Java stack (the run time stack). Each *frame* of the Java stack contains a *local variable array* and an *operand stack*.
- heap (contains data associated with objects). Periodically, *garbage collection* deallocates objects in the heap that are no longer referenced.

There are two types of stacks in the JVM

- The *Java stack*
- The Java stack consists of frames, one frame for each method invocation. Each frame contains an *operand stack* and a local variable array.

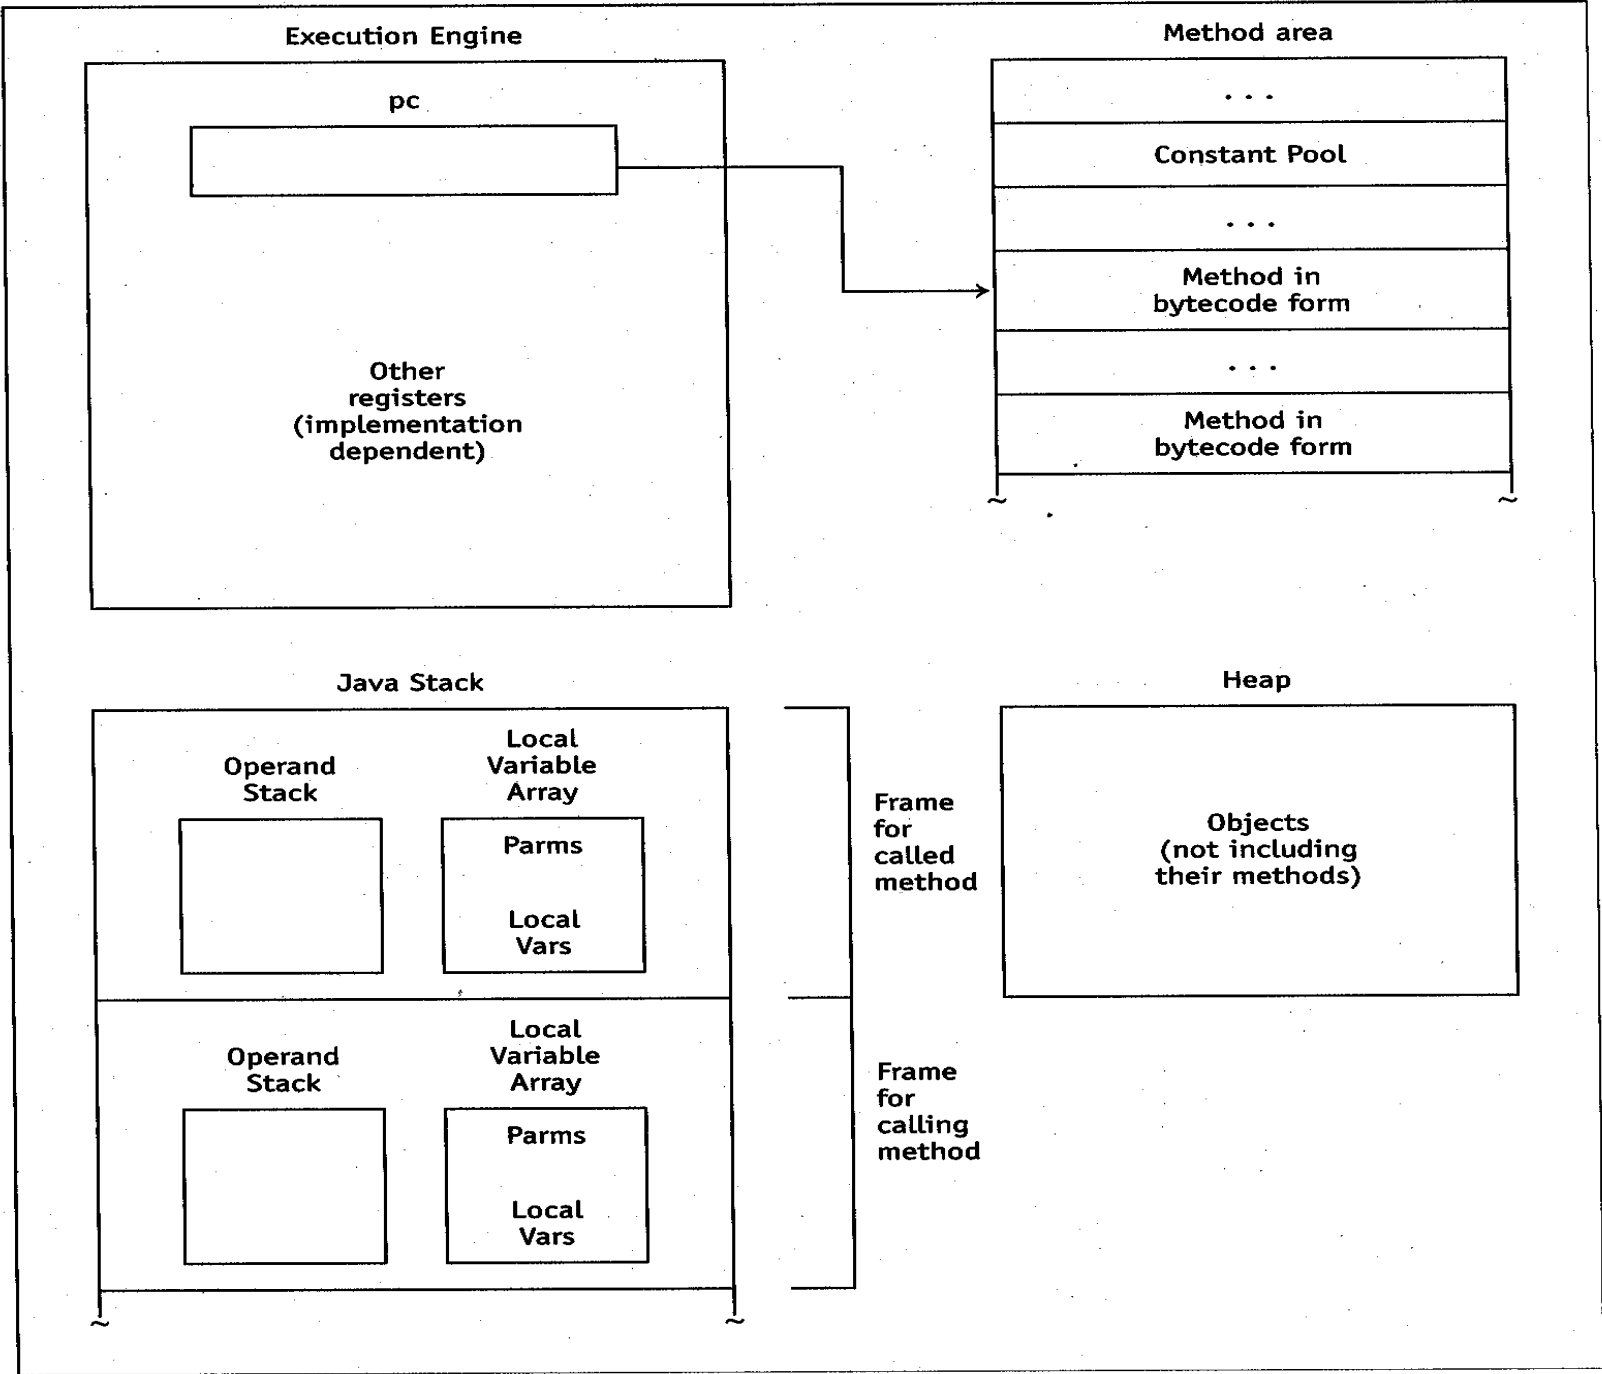
Local variable array

Contains local variables numbered starting from 0. For example, the first slot of the local variable array is called local variable 0.

Operand stack

Used to hold operands and results during the execution of instructions.

FIGURE 16.1



Some instructions consist of an opcode only. For example,

iconst_0, iconst_1, iconst_2,
iconst_3, iconst_4, iconst_5

which push 0, 1, 2, 3, 4, and 5,
respectively, onto the operand stack.

The more common operations are
performed by such opcode-only
instructions.

Some instructions require an operand.
For example,

bipush 6

which pushes 6. This instruction consists of the opcode for bipush followed by a byte containing 6.

To push a number greater than 127, use sipush (short int push). For example,

sipush 130

Symbolic bytecode that adds three numbers

Let's say we wish to add 3, 6, and 130. One possible instruction sequence is

```
iconst_3      ; push 3
bipush 6      ; push 6
iadd          ; pop 6 and 3, add, and push sum (9)
sipush 130    ; push 130
iadd          ; add 130 and 9, add, and push sum
```

The initial letter of some mnemonics indicates the data type. For example, iadd, dadd, fadd, ladd.

a: reference

d: double

f: float

i: integer

ia: integer array

l: long

Load instructions on the JVM

`iload_0`

pushes the value in local variable 0 (i.e., it pushes the value from the first slot of the local variable array onto the operand stack.)

`iload 4`

pushes the value in local variable 4.

Store instructions on the JVM

`istore_0`

pops and stores the value on top of the operand stack into local variable 0.

`istore 4`

pops and stores the value on top of the operand stack into local variable 4.

A static variable in Java is a variable associated with a class rather than an object. It is shared by all objects of its class.

A static method in Java is a method that can be called via its class.

The `getstatic` and `putstatic` instructions transfer values between the top of the operand stack and static variables.

The operand that appears in `getstatic` and `putstatic` instructions is an index into the *constant pool*. For example,

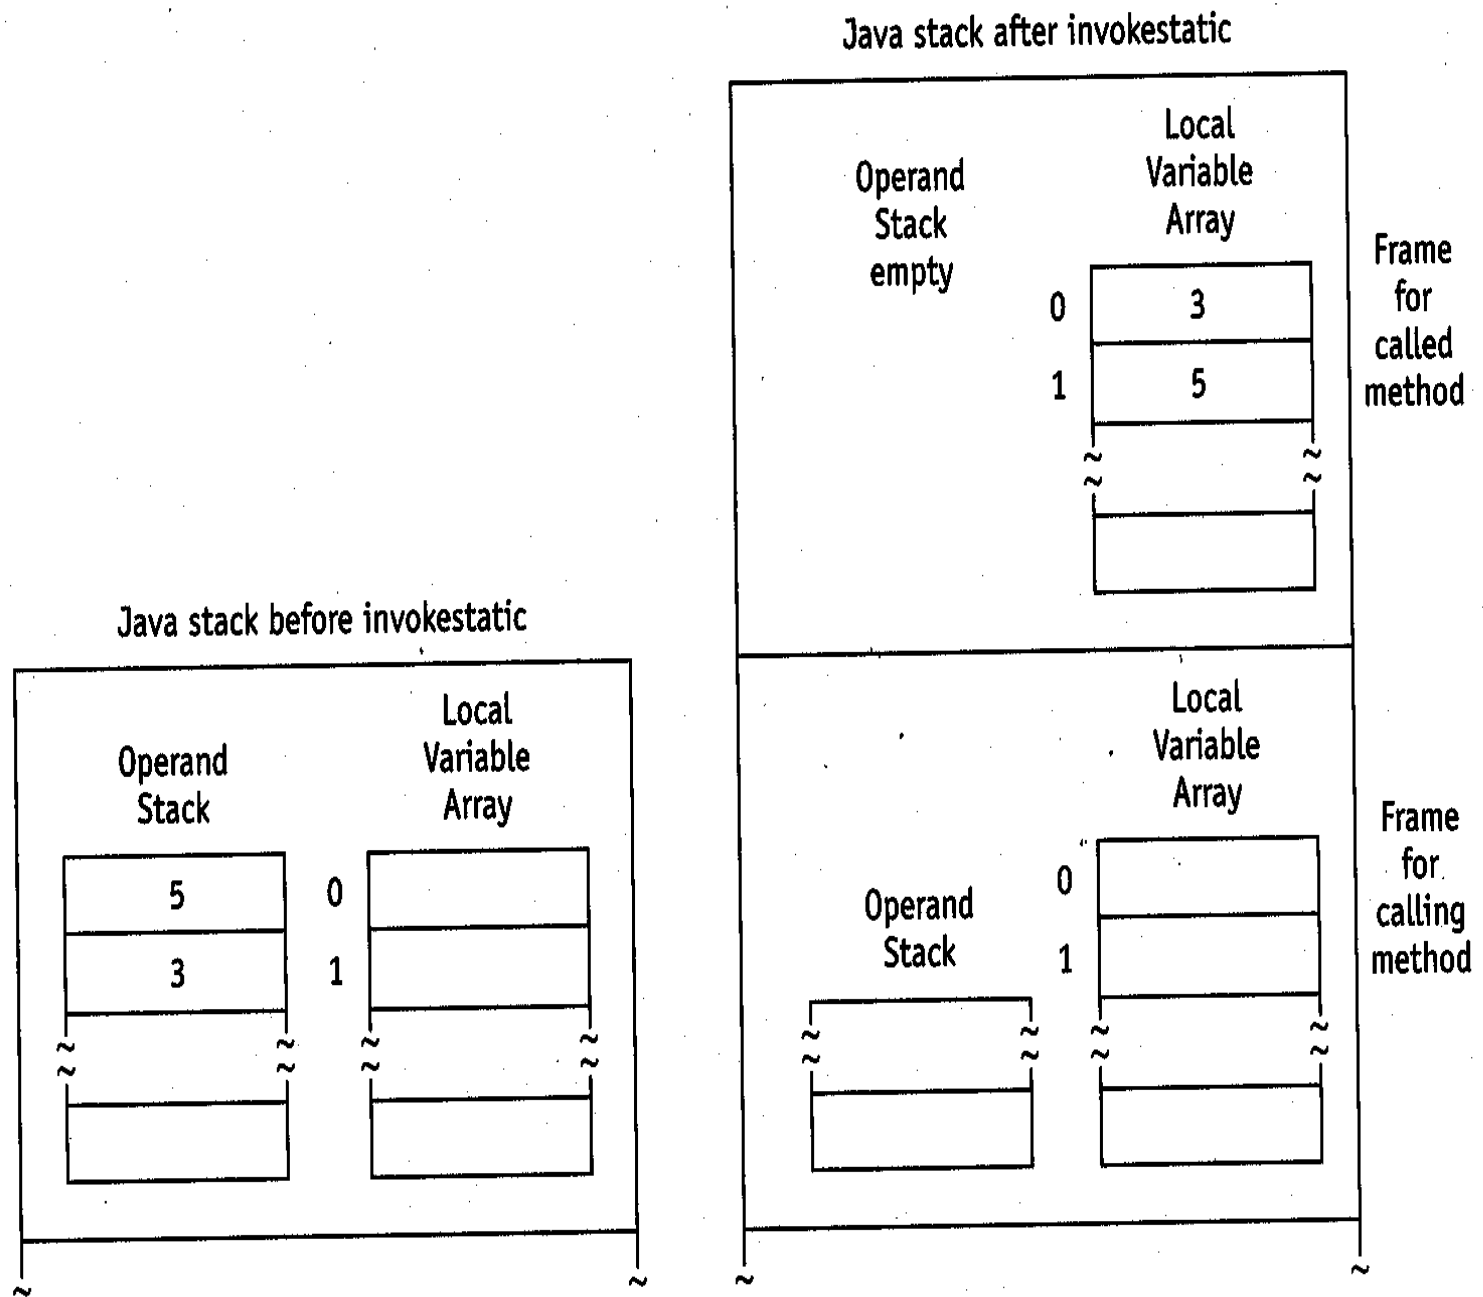
`getstatic 2`

2 is a constant pool index.

Invoking a static method with invokestatic instruction

- Creates frame for the called method and pushes it onto the Java stack.
- Pops the arguments from the caller's operand stack and places them in the called method's local variable array starting from local variable 0.
- Transfers control to the called method.

FIGURE 16.2



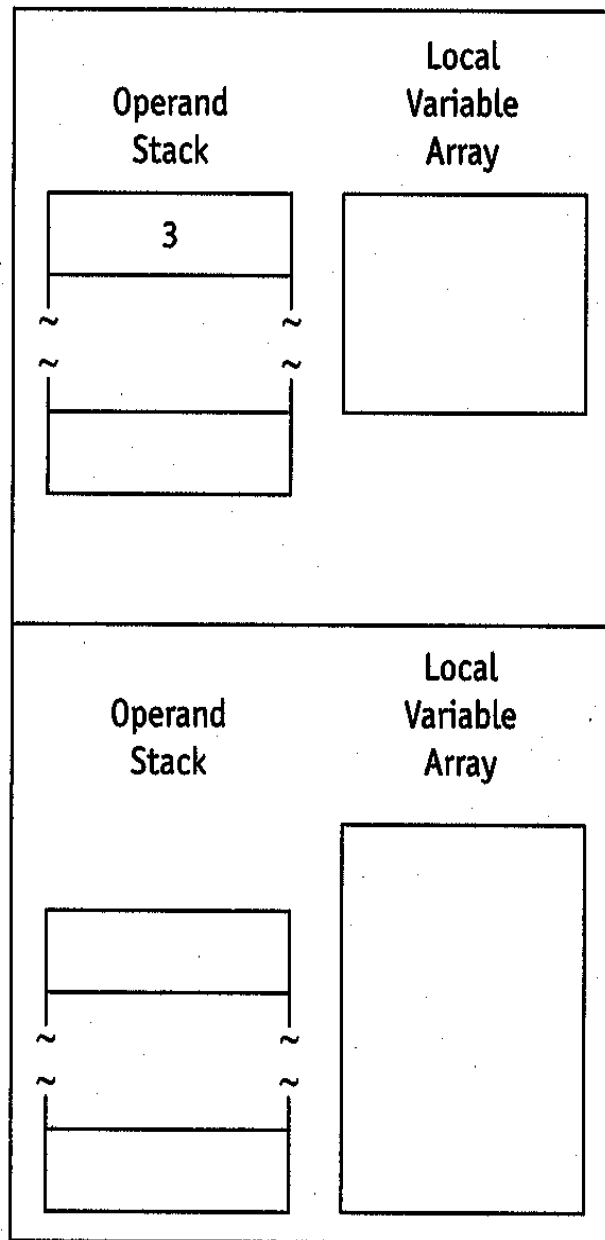
Returning a value to the calling method with the ireturn instruction

```
icont_3      ; push 3  
ireturn      ; return 3
```

The value returned is pushed onto the calling method's operand stack.


FIGURE 16.3

Java stack before ireturn



Frame
for called
method

Java stack after ireturn



Frame
for calling
method

Implementation of the execution engine

The execution engine of the JVM repeatedly performs the four steps that a CPU typically performs:

1. Fetch the instruction.
2. Increment the `pc`.
3. Decode the opcode.
4. Execute the instruction.

FIGURE 16.4

Action

C++ Code

```

1 Fetch          | opcode =
2                |   (unsigned char)*pc;   // zero extend *pc
3                |
4 Increment pc   | pc++;                  // move pc over opcode
5                |
6 Decode         | switch switch (opcode) {
7 Opcode         |
8                |   case 16:              // bipush opcode is 16
9                |       sp--;             // make room for push
10               |       sp[0] = (int)*pc;  // fetch and push operand
11               |       pc++;             // move pc over operand
12               |       break;
13               |
14               |   case 27:              // iload_1 opcode is 27
15               |       sp--;             // make room for push
16               |       sp[0] = ap[1];     // push loc var 1
17               |       break;
18 Execute       |
19               |   case 96:              // iadd opcode is 96
20               |       sp[1] { =         // slot below top
21               |       sp[0] + sp[1];    // sum of top 2 slots
22               |       sp++;             // now pop top
23               |       break;
24               |
25               |   // other cases
26               | }

```


The wisdom of using a stack architecture for the JVM

- A stack architecture on a simulated machine is no slower than a register architecture.
- Bytecode is very compact which is important for a web programs.

A simple Java program follows,
along with its bytecode

FIGURE 16.5

```

1 class Simple {
2   static int gv1, gv2 = 5;
3
4   // <init> method                ; default constructor
5   // 0  2A      aload_0            ; get object's reference
6   // 1  B70001 invokespecial 1     ; invoke <init> in superclass
7   // 4  B1      return
8   //=====
9   public static void main(String arg[])
10  {
11    int lv1,
12    lv2 = 7;      // 0  1007  bipush 7      ; push 7
13                  // 2  3D    istore_2      ; store in lv2
14
15    lv1 = 11;     // 3  100B  bipush 11     ; push 11
16                  // 5  3C    istore_1      ; store in lv1
17
18    gv1 = fa(gv2, lv1, lv2);
19                  // 6  B20002 getstatic 2   ; push gv2
20                  // 9  1B     iload_1      ; push lv1
21                  // 10 1C     iload_2      ; push lv2
22                  // 11 B80003 invokestatic 3 ; call fa
23                  // 14 B30004 putstatic 4   ; pop into gv1
24                  // 17 B1     return
25  }

```

(continued)

FIGURE 16.5 (continued)

```
26 //=====
27 public static int fa(int x, int y, int z)
28 {
29     return x + y + z; // 0  1A  iload_0      ; push x
30                       // 1  1B  iload_1      ; push y
31                       // 2  60  iadd         ; pop/pop/add/push
32                       // 3  1C  iload_2      ; push z
33                       // 4  60  iadd         ; pop/pop/add/push
34                       // 5  AC  ireturn      ; pop and return
35 }
36 //=====
37 // <clinit> method          ; class initializer
38 // 0  08  iconst_5           ; push 5
39 // 1  B30002 putstatic 2      ; pop into gv2
40 // 4  B1  return
41 }
```

A formatted display of the
constant pool for our simple
program follows.

FIGURE 16.6

3 (operand in the invokestatic instruction on line 22 in Fig. 16.5)

Constant Pool					
Index	Tag	Information			
1:	10 (Methodref)	6	(Class Index);	21	(NameAndType Index)
2:	9 (Fieldref)	5	(Class Index);	22	(NameAndType Index)
→ 3:	10 (Methodref)	5	(Class Index);	23	(NameAndType Index)
4:	9 (Fieldref)	5	(Class Index);	24	(NameAndType Index)
→ 5:	7 (Class)	25			
6:	7 (Class)	26			
7:	1 (UTF8)	"gvl"			
8:	1 (UTF8)	"I"			
9:	1 (UTF8)	"gv2"			
10:	1 (UTF8)	"<init>"			
11:	1 (UTF8)	"()V"			
12:	1 (UTF8)	"Code"			
13:	1 (UTF8)	"LineNumberTable"			
14:	1 (UTF8)	"main"			
15:	1 (UTF8)	"([Ljava/lang/string;)V"			
16:	1 (UTF8)	"fa"			
17:	1 (UTF8)	"(III)I"			
18:	1 (UTF8)	"<clinit>"			
19:	1 (UTF8)	"SourceFile"			
20:	1 (UTF8)	"Simple.java"			
21:	12 (NameAndType)	10 (Name Index);	11	(Descriptor Index)	
22:	12 (NameAndType)	9 (Name Index);	8	(Descriptor Index)	
23:	12 (NameAndType)	16 (Name Index);	17	(Descriptor Index)	←
24:	12 (NameAndType)	7 (Name Index);	8	(Descriptor Index)	
→ 25:	1 (UTF8)	"Simple"			
26:	1 (UTF8)	"java/lang/Object"			

Information in the constant pool for index 3

3 -> 5 -> 25 "Simple"

which yields "Simple", the class name of the method. Similarly, the chains

3 -> 23 -> 16 "fa"

3 -> 23 -> 17 "(III)I"

An attribute in a class file

```
0013      Attribute name index ("SourceFile")
00000002  Attribute length (length of what follows this field)
0014      Name index ("Simple.java")
```

The first entry is the constant pool index of the attribute name. The second entry is the length of what follows.

A hex display of the complete class file for our simple program follows.

FIGURE 16.7

Hex Display of Simple.class

```
1      CAFE BABE  Magic number (signature for class files)
2      0003      Minor version number of JVM
3      002D      Major version number of JVM
4  -----
5 Constant pool
6
7      001B      Constant pool count
8 Index  Tag
9 01:    0A 0006 0015
10 02:    09 0005 0016
11 03:    0A 0005 0017
12 04:    09 0005 0018
13 05:    07 0019
14 06:    07 001A
15 07:    01 0003 6776 31      "gv1"
16 08:    01 0001 49          "I"
17 09:    01 0003 6776 32      "gv2"
18 0A:    01 0006 3C69 6E69 743E  "<init>"
19 0B:    01 0003 2829 56      "()V"
```

(continued)

FIGURE 16.7 (continued)

```

20 0C:      01  0004 436F 6465                                "Code"
21 0D:      01  000F 4C69 6E65 4E75 6D62 6572 5461 626C 65
22                                                "LineNumberTable"
23 0E:      01  0004 6D61 696E                                "main"
24 0F:      01  0016 285B 4C6A 6176 612F 6C61 6E67 2F53 7472 696E 673B 2956
25                                                "([Ljava/lang/String;)V"
26 10:      01  0002 6661                                "fa"
27 11:      01  0006 2849 4949 2949                                "(III)I"
28 12:      01  0008 3C63 6C69 6E69 743E                                "<clinit>"
29 13:      01  000A 536F 7572 6365 4669 6C65                                "SourceFile"
30 14:      01  000B 53 696D 706C 652E 6A61 7661                                "Simple.java"
31 15:      0C  000A 000B
32 16:      0C  0009 0008
33 17:      0C  0010 0011
34 18:      0C  0007 0008
35 19:      01  0006 5369 6D70 6C65                                "Simple"
36 1A:      01  0010 6A61 7661 2F6C 616E 672F 4F62 6A65 6374
37                                                "java/lang/Object"
38 -----
39 0020      Access flags (0020 indicates that JVM should
40           use newer version of invokespecial instruction)
41 0005      This class index ("Simple")
42 0006      Super class index ("java/lang/Object")
43 0000      Interfaces count
44 0002      Fields count
45 -----
46 gv1
47 0008      Access flags (0008 indicates static)
48 0007      Name index ("gv1")
49 0008      Descriptor index ("I")
50 0000      Attributes count
51 -----
52 gv2
53 0008      Access flags (0008 indicates static)
54 0009      Name index ("gv2")
55 0008      Descriptor index ("I")
56 0000      Attributes count
57 -----
58 0004      Number of methods
59 -----
60 <init>
61 0000      Access flags (0000 indicates default access)
62 000A      Name index of method ("<init>")
63 000B      Descriptor index      ("()V")
64 0001      Attributes count

```

(continued)

65				
66	000C	Attribute name index ("Code")		
67	0000001D	Attribute length		
68	0001	Max stack		
69	0001	Max locals		
70	00000005	Code length		
71	2A B70001 B1	Bytecode		Code
72	0000	Exceptions count		Attribute
73	0001	Attributes count		
74				
75	000D	Attribute Name Index ("LineNumberTable")		Line
76	00000006	Attribute length		Number
77	0001	LineNumberTable length		Table
78	0000 0001	Location/line number		Attribute
79				
80	-----			
81	main			
82	0009	Access flags (0009 indicates public and static)		
83	000E	Name index of method ("main")		
84	000F	Descriptor index ("([Ljava/lang/String;)V")		
85	0001	Attribute count		
86				
87	000C	Attribute name index ("Code")		
88	00000036	Attribute length		
89	0003	Max stack		
90	0003	Max locals		
91	00000012	Code length		
92	1007 3D 100B 3C B20002 1B 1C B80003 B30004 B1	Bytecode		Code
93	0000	Exceptions count		Attribute
94	0001	Attribute count		
95				
96	000D	Attribute name index ("LineNumberTable")		Line
97	00000012	Attribute length		Number
98	0004	LineNumberTable length		Table
99	0000 000C	Location/line number		Attribute
100	0003 000F	Location/line number		
101	0006 0012	Location/line number		
102	0011 0019	Location/line number		
103				
104	-----			
105	fa			
106	0009	Access flags (0009 indicates public and static)		
107	0010	Name index of method ("fa")		
108	0011	Descriptor index ("(III)I")		
109	0001	Attribute count		

(continued)

FIGURE 16.7 (continued)

110							
111	000C	Attribute name index ("Code")					
112	0000001E	Attribute length					
113	0002	Max stack					
114	0003	Max locals					
115	00000006	Code length					
116	1A 1B 60 1C 60 AC	Bytecode					
117	0000	Exceptions count					
118	0001	Attribute count					
119							
120	000D	Attribute name index ("LineNumberTable")					
121	00000006	Attribute length					
122	0001	LineNumberTable length					
123	0000 001D	Location/line number					
124							
125	-----						
126	<clinit>						
127	0008	Access flags (0008 indicates static)					
128	0012	Name index of method ("<clinit>")					
129	000B	Descriptor index ("()V")					
130	0001	Attribute count					
131							
132	000C	Attribute name index ("Code")					
133	0000001D	Attribute length					
134	0001	Max stack					
135	0000	Max locals					
136	00000005	Code length					
137	08 B30002 B1	Bytecode					
138	0000	Exceptions count					
139	0001	Attribute count					
140							
141	000D	Attribute name index ("LineNumberTable")					
142	00000006	Attribute length					
143	0001	LineNumberTable length					
144	0000 0002	Location/line number					
145							
146	-----						
147	0001	Attributes count					
148							
149	0013	Attribute name index ("SourceFile")					
150	00000002	Attribute length					
151	0014	Name index ("Simple.java")					

Sizes of comparable programs

FIGURE 16.8

	SPARC	Pentium	Bytecode
main	40	52	18
fa	12	14	6
<init>	-	-	5
<clinit>	-	-	5
Total bytes	52	66	34

Some comparison and control instructions

- goto unconditional jump
 - if_cmplt compares top two stack items
 - if_icmpge compares top two stack items
 - iflt compares top of stack with 0
 - if_acmpeq compares references
 - if_acmpne compares references
- See the illustrative program on the next slide.

FIGURE 16.9

```

1 class Control {
2   public static void main(String arg[])
3   {
4       int x = 3,           // 0    06          iconst_3   ; push 3
5                           // 1    3C          istore_1   ; pop into x
6
7       y = 4;              // 2    07          iconst_4   ; push 4
8                           // 3    3D          istore_2   ; pop into y
9
10      while (x < 10 )      // 4    A7 0006      goto 6      ; goto loc 4+6
11
12          x++;             // 7    84 01 01      iinc 1 1    ; add 1 to 1
13
14      // (x < 10)           10   1B          iload_1    ; push x
15      // exit test         11   10 0A        bipush 10   ; push 10
16      // is here           13   A1 FFFA      if_icmplt -6 ; goto loc 13-6
17
18      if (x < y)            // 16   1B          iload_1    ; push x
19                           // 17   1C          iload_2    ; push y
20                           // 18   A2 0009      if_icmpge 9 ; goto 18+9
21
22          x = 20;          // 21   10 14        bipush 20   ; push 20
23                           // 23   3C          istore_1    ; store in x
24
25      else                 // 24   A7 0006      goto 6      ; goto 24+6

```

(continued)

FIGURE 16.9 (continued)

```
26
27      x = 30;      // 27 10 1E      bipush 30 ; push 30
28                        // 29 3C      istore_1 ; pop into x
29
30      }            // 30 B1      return ; return to caller
31 }
```

Instructions that jump use pc-relative addressing

A70006 (the machine code for goto 6) jumps to the location whose address is $6 +$ the contents of the pc register (before incrementation).

Unassembling the Simple class file

```
javap -c Simple
```

FIGURE 16.10

```
1 Compiled from Simple.java
2 class Simple extends java.lang.Object {
3     static int gv1;
4     static int gv2;
5     Simple();
6     public static void main(java.lang.String[]);
7     public static int fa(int, int, int);
8     static {};
9 }
10
11 Method Simple()          ← <init> method
12     0 aload_0
13     1 invokespecial #1 <Method java.lang.Object()>
14     4 return
15
16 Method void main(java.lang.String[])
17     0 bipush 7             constant pool index
18     2 istore_2
19     3 bipush 11
20     5 istore_1             symbolic info obtained via constant
                             pool index 2
21     6 getstatic #2 <Field int gv2>
22     9 iload_1
23     10 iload_2
24     11 invokestatic #3 <Method int fa(int, int, int)>
25     14 putstatic #4 <Field int gv1>
26     17 return
27
28 Method int fa(int, int, int)
29     0 iload_0
30     1 iload_1
31     2 iadd
32     3 iload_2
33     4 iadd
34     5 ireturn
35
36 Method static {}        ← <clinit> method
37     0 iconst_5
38     1 putstatic #2 <Field int gv2>
39     4 return
```

Unassemble this program to see its bytecode

FIGURE 16.11 a)

```
1 class IRTest {
2     public static void main(String arg[])
3     {
4         int y;
5         y = sum(1, 2);
6     }
7     static int sum(int m, int n)
8     {
9         return m + n;
10    }
11 }
```

```

1 Method void main(java.lang.String[])
2     0 iconst_1          ; push 1
3     1 iconst_2          ; push 2
4     2 invokestatic #2 <Method int sum(int, int)>
5     5 istore_1          ; save value returned on stack in loc var 1
6     6 return
7
8 Method int sum(int, int)
           ; return m + n;
9     0 iload_0           ; get 1st parameter from loc var 0
10    1 iload_1           ; get 2nd parameter from loc var 1
11    2 iadd              ; pop/pop/add/push
12    3 ireturn           ; return value on top of stack

```

Arrays and objects

FIGURE 16.12 a)

```
1 class OATest {
2     public static void main(String arg[])
3     {
4         O o;
5         int a[];
6         o = new O();
7         o.x = 6;
8         o.f();
9         a = new int[7];
10        a[5] = 3;
11    }
12 }
```

b)

```
1 Method void main(java.lang.String[])
2     0 new #2 <Class O> ; o = new O();
3     3 dup ; duplicate reference
4     4 invokespecial #3 <Method O()>
5     7 astore_1 ; save reference
6
7     8 aload_1 ; o.x = 6;
8     9 bipush 6
9    11 putfield #4 <Field int x>
10
11    14 aload_1 ; o.f();
12    15 invokevirtual #5 <Method void f()>
13
14    18 bipush 7 ; a = new int[7];
15    20 newarray int
16    22 astore_2 ; save reference
17
18    23 aload_2 ; a[5] = 3;
19    24 iconst_5
20    25 iconst_3
21    26 iastore
22
23    27 return
```


Now that you know the basics of the JVM,
you can enjoy (and understand) some more
advanced discussions of the JVM.

FURTHER READING



Engel, J. *Programming for the Java Virtual Machine*. Reading, MA: Addison-Wesley, 1999.

Harold, E. *Java Secrets*. Forster City, CA: IDG Books Worldwide, 1997.

Lindholm, T. and Yellin, F. *The Java Virtual Machine Specification*. Reading, MA: Addison-Wesley, 1997.

Meyer, J. and Downing, T. *Java Virtual Machine*. Sebastopol, CA: O'Reilly, 1997.

Venners, B. *Inside the Java Virtual Machine*. New York: McGraw-Hill, 1998.

return 0;